

# AQI\_Data\_Cleaning

October 19, 2024

## #Importing Necessary Libraries and Loading Data

```
[1]: # Importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.impute import KNNImputer
```

```
[2]: # Load the Data
file_path = '/content/air_pollution_data.csv'
df = pd.read_csv(file_path)
```

## #Data Exploration

```
[3]: # Display first few rows
print("Initial Data Snapshot:")
print(df.head())
```

Initial Data Snapshot:

|   | city      | date       | aqi | co      | no    | no2   | o3     | so2   | pm2_5  | \ |
|---|-----------|------------|-----|---------|-------|-------|--------|-------|--------|---|
| 0 | Ahmedabad | 30-11-2020 | 5   | 520.71  | 2.38  | 16.28 | 130.18 | 47.68 | 65.96  |   |
| 1 | Ahmedabad | 01-12-2020 | 5   | 1682.28 | 7.71  | 54.84 | 0.73   | 21.70 | 120.95 |   |
| 2 | Ahmedabad | 02-12-2020 | 5   | 1815.80 | 16.54 | 49.35 | 0.17   | 23.84 | 133.47 |   |
| 3 | Ahmedabad | 03-12-2020 | 5   | 2296.45 | 41.57 | 40.10 | 0.00   | 35.76 | 150.37 |   |
| 4 | Ahmedabad | 04-12-2020 | 5   | 2189.64 | 23.92 | 58.95 | 0.02   | 28.13 | 160.79 |   |

|   | pm10   | nh3   |
|---|--------|-------|
| 0 | 72.13  | 8.36  |
| 1 | 154.53 | 27.36 |
| 2 | 172.63 | 28.12 |
| 3 | 202.15 | 36.48 |
| 4 | 205.80 | 40.53 |

```
[5]: # Initial Data Exploration
print("\nData Information:")
print(df.info())
```

Data Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 23504 entries, 0 to 23503

Data columns (total 11 columns):

| #  | Column | Non-Null Count | Dtype   |
|----|--------|----------------|---------|
| 0  | city   | 23504 non-null | object  |
| 1  | date   | 23504 non-null | object  |
| 2  | aqi    | 23504 non-null | int64   |
| 3  | co     | 23504 non-null | float64 |
| 4  | no     | 23504 non-null | float64 |
| 5  | no2    | 23504 non-null | float64 |
| 6  | o3     | 23504 non-null | float64 |
| 7  | so2    | 23504 non-null | float64 |
| 8  | pm2_5  | 23504 non-null | float64 |
| 9  | pm10   | 23504 non-null | float64 |
| 10 | nh3    | 23504 non-null | float64 |

dtypes: float64(8), int64(1), object(2)

memory usage: 2.0+ MB

None

```
[6]: print("\nStatistical Summary:")
      print(df.describe())
```

Statistical Summary:

|       | aqi          | co           | no           | no2          | o3 \         |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 23504.000000 | 23504.000000 | 23504.000000 | 23504.000000 | 23504.000000 |
| mean  | 3.920354     | 1113.224543  | 6.00554      | 25.044104    | 35.059777    |
| std   | 1.415490     | 1401.770372  | 24.50272     | 25.839242    | 31.901760    |
| min   | 1.000000     | 173.570000   | 0.00000      | 0.310000     | 0.000000     |
| 25%   | 3.000000     | 447.270000   | 0.00000      | 8.740000     | 7.870000     |
| 50%   | 5.000000     | 700.950000   | 0.00000      | 16.450000    | 28.250000    |
| 75%   | 5.000000     | 1188.280000  | 0.27000      | 32.220000    | 54.360000    |
| max   | 5.000000     | 23071.290000 | 457.76000    | 331.760000   | 406.270000   |

|       | so2          | pm2_5        | pm10         | nh3          |
|-------|--------------|--------------|--------------|--------------|
| count | 23504.000000 | 23504.000000 | 23504.000000 | 23504.000000 |
| mean  | 15.971449    | 98.598310    | 121.848091   | 12.060212    |
| std   | 23.943464    | 135.572391   | 160.429589   | 17.544759    |
| min   | 0.190000     | 0.500000     | 0.580000     | 0.000000     |
| 25%   | 4.470000     | 24.677500    | 32.277500    | 2.340000     |
| 50%   | 7.990000     | 58.860000    | 75.775000    | 6.520000     |
| 75%   | 16.450000    | 117.605000   | 147.642500   | 15.830000    |
| max   | 442.510000   | 2203.550000  | 2429.130000  | 352.620000   |

#Data Cleaning

## #1) Handling Missing Values Using KNNImputer

```
[7]: # List of pollutant columns where we want to treat 0 as missing
pollutant_cols = ['co', 'no', 'no2', 'o3', 'so2', 'pm2_5', 'pm10', 'nh3']
```

```
[9]: # Print the number of 0s in each column before cleaning
print("\nNumber of 0s in each column before cleaning:")
for col in pollutant_cols:
    num_zero = (df[col] == 0).sum()
    print(f"{col}: {num_zero}")
```

Number of 0s in each column before cleaning:

```
co: 0
no: 12740
no2: 0
o3: 1551
so2: 0
pm2_5: 0
pm10: 0
nh3: 223
```

```
[10]: # Initialize KNNImputer (0 values will be treated as missing automatically by
      ↪ setting missing_values=0)
imputer = KNNImputer(n_neighbors=5, missing_values=0)
```

```
[11]: # Apply imputer on the pollutant columns
df_imputed = imputer.fit_transform(df[pollutant_cols])

# Replace the original columns with the imputed values
df[pollutant_cols] = df_imputed
```

```
[12]: # Print the number of 0s in each column after cleaning
print("\nNumber of 0s in each column after cleaning:")
for col in pollutant_cols:
    num_zero = (df[col] == 0).sum()
    print(f"{col}: {num_zero}")
```

Number of 0s in each column after cleaning:

```
co: 0
no: 0
no2: 0
o3: 0
so2: 0
pm2_5: 0
pm10: 0
nh3: 0
```

## #Data Transformation for Dashboard

```
[13]: # Convert 'date' column to datetime format and extract 'year' and 'month'
df['date'] = pd.to_datetime(df['date'], format='%d-%m-%Y')
df['year'] = df['date'].dt.year
df['month'] = df['date'].dt.month
```

```
[14]: # Check for duplicate rows
duplicates = df.duplicated().sum()
if duplicates:
    print(f"Found {duplicates} duplicate rows. Dropping them.")
    df = df.drop_duplicates()
else:
    print("No duplicates found.")
```

No duplicates found.

## #Visualization with Matplotlib and Seaborn

```
[16]: # City-wise Data Count
city_counts = df['city'].value_counts()

# Convert city counts to a DataFrame for a cleaner table display
city_counts_df = city_counts.reset_index()
city_counts_df.columns = ['City', 'Data Count']

# Print the city-wise data count as a table
print("\nCity-wise Data Count:")
print(city_counts_df)
```

City-wise Data Count:

|    | City               | Data Count |
|----|--------------------|------------|
| 0  | Ahmedabad          | 904        |
| 1  | Aizawl             | 904        |
| 2  | Thiruvananthapuram | 904        |
| 3  | Talcher            | 904        |
| 4  | Shillong           | 904        |
| 5  | Patna              | 904        |
| 6  | Mumbai             | 904        |
| 7  | Lucknow            | 904        |
| 8  | Kolkata            | 904        |
| 9  | Kochi              | 904        |
| 10 | Jorapokhar         | 904        |
| 11 | Jaipur             | 904        |
| 12 | Hyderabad          | 904        |
| 13 | Guwahati           | 904        |
| 14 | Gurugram           | 904        |
| 15 | Ernakulam          | 904        |

|    |               |     |
|----|---------------|-----|
| 16 | Delhi         | 904 |
| 17 | Coimbatore    | 904 |
| 18 | Chennai       | 904 |
| 19 | Chandigarh    | 904 |
| 20 | Brajrajnagar  | 904 |
| 21 | Bhopal        | 904 |
| 22 | Bengaluru     | 904 |
| 23 | Amritsar      | 904 |
| 24 | Amaravati     | 904 |
| 25 | Visakhapatnam | 904 |

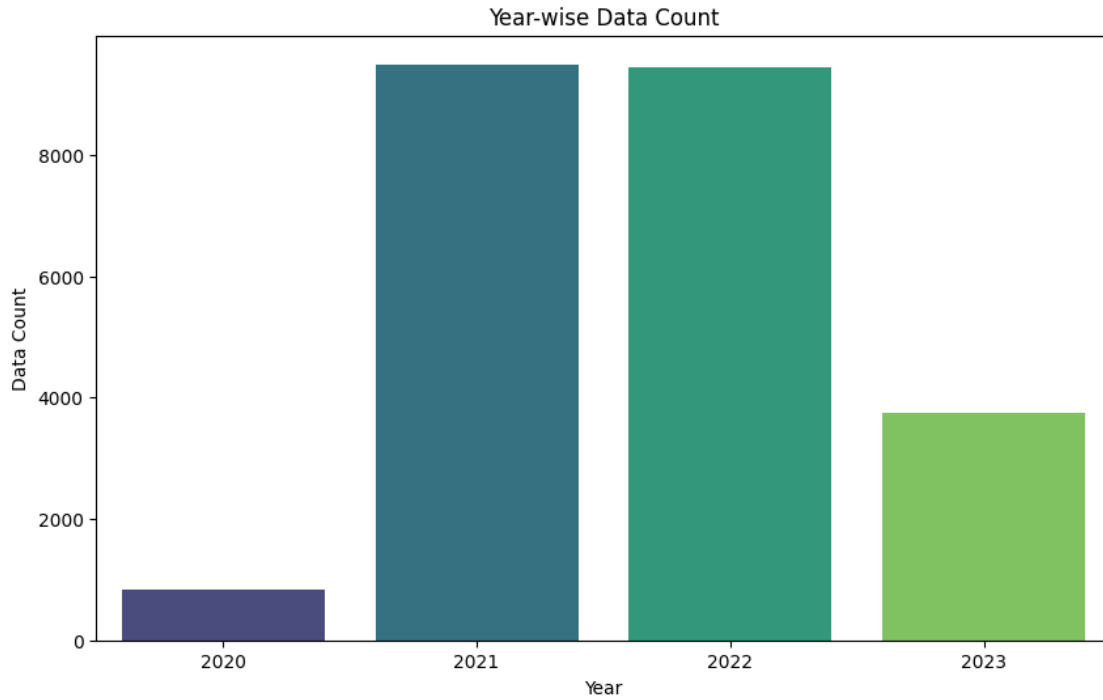
```
[17]: # Year-wise Data Count
year_counts = df['year'].value_counts().sort_index()

# Plot Year-wise Data Count
plt.figure(figsize=(10,6))
sns.barplot(x=year_counts.index, y=year_counts.values, palette='viridis')
plt.title('Year-wise Data Count')
plt.xlabel('Year')
plt.ylabel('Data Count')
plt.show()
```

<ipython-input-17-4982ac4a5279>:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=year_counts.index, y=year_counts.values, palette='viridis')
```



[19]: *# Detect and Visualize Outliers Using Boxplots*

```
def detect_outliers_iqr(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
    return outliers
```

[21]: columns\_to\_check = ['aqi', 'co', 'no', 'no2', 'o3', 'so2', 'pm2\_5', 'pm10',  
↪ 'nh3']

[22]: *# Detect outliers for each column and print the count*

```
for col in columns_to_check:
    outliers = detect_outliers_iqr(df, col)
    print(f"Number of outliers in '{col}': {len(outliers)}")
```

```
Number of outliers in 'aqi': 0
Number of outliers in 'co': 2257
Number of outliers in 'no': 3558
Number of outliers in 'no2': 1515
Number of outliers in 'o3': 283
Number of outliers in 'so2': 2694
```

Number of outliers in 'pm2\_5': 1850  
Number of outliers in 'pm10': 1746  
Number of outliers in 'nh3': 1394

```
[23]: # Plot Boxplots for Outliers Detection
plt.figure(figsize=(12,8))
for i, col in enumerate(columns_to_check, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(y=df[col], palette="Set2")
    plt.title(f'Boxplot for {col}')
plt.tight_layout()
plt.show()
```

<ipython-input-23-209e5b9b6d71>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

<ipython-input-23-209e5b9b6d71>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

<ipython-input-23-209e5b9b6d71>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

<ipython-input-23-209e5b9b6d71>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

<ipython-input-23-209e5b9b6d71>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

```
<ipython-input-23-209e5b9b6d71>:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

```
<ipython-input-23-209e5b9b6d71>:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```

```
<ipython-input-23-209e5b9b6d71>:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

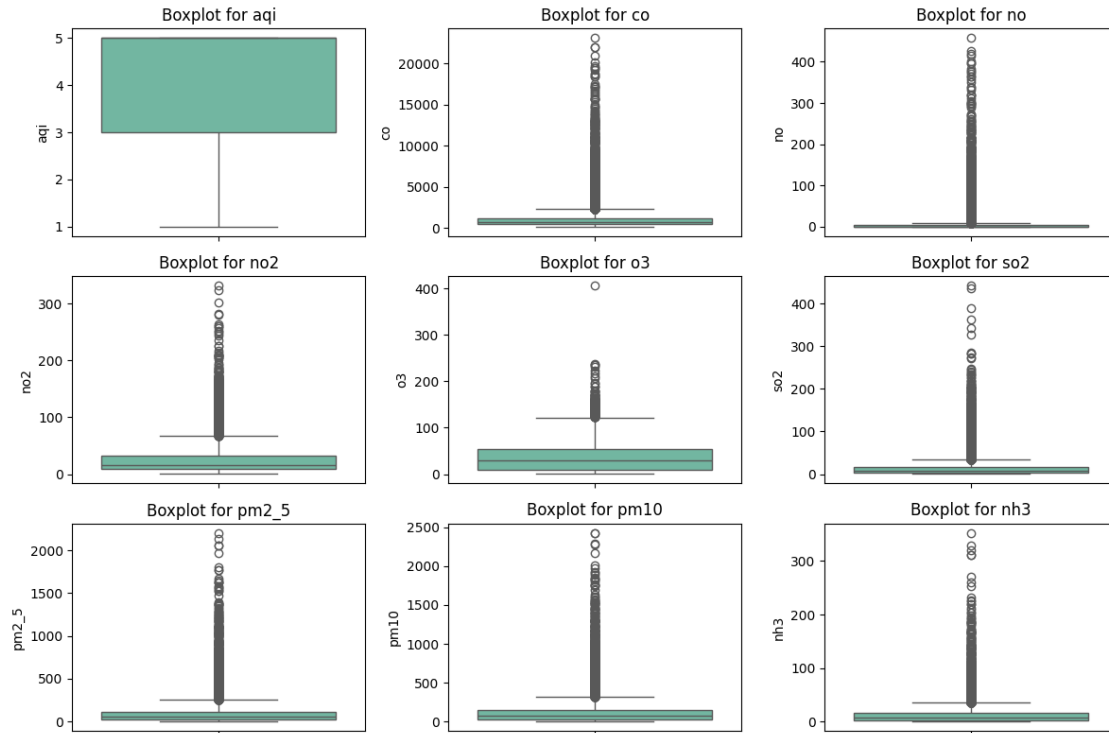
```
sns.boxplot(y=df[col], palette="Set2")
```

```
<ipython-input-23-209e5b9b6d71>:5: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(y=df[col], palette="Set2")
```





### #Why Handling Outliers May Not Be Necessary for Power BI Dashboard?

Handling outliers isn't always necessary for Power BI dashboards since their purpose is to visualize data trends, and outliers can provide valuable insights, especially in environmental data like air pollution. Removing them could result in losing important information about extreme events. Power BI can display outliers clearly, allowing users to analyze their impact without altering the dataset. Thus, it's often better to retain outliers for transparency.

```
[24]: # Save the cleaned data to a CSV file
cleaned_file_path = '/content/air_pollution_cleaned_data.csv'
df.to_csv(cleaned_file_path, index=False)
print(f"\nCleaned data saved to {cleaned_file_path}")
```

Cleaned data saved to /content/air\_pollution\_cleaned\_data.csv